



University
of Glasgow | School of
Computing Science

A Deep Learning Approach to Musical Effects

Kieran McCool

School of Computing Science
Sir Alwyn Williams Building
University of Glasgow
G12 8QQ

Level 4 Project — March 28, 2018

Contents

1	Introduction	2
1.1	Musical Effects	2
1.2	Musical Effect Modelling	2
1.3	Machine Learning	3
2	Background	4
2.1	Nebula VST	4
2.2	Project Magenta	4
2.3	WaveNet	4
3	Aims	6
3.1	Effect Choices	6
3.2	Deep Learning	6
4	Methods	8
4.1	Network Architecture	8
4.1.1	Convolutional Networks	8
4.1.2	Long-Short Term Networks	8
4.1.3	WaveNet	8
4.2	Evaluating Success	10
4.2.1	Melspectograms	10
4.2.2	Impulse Response	10
4.2.3	ABX Testing	10

5	Implementation	13
5.1	Generating Test Data	13
5.1.1	Random Audio Signals	13
5.1.2	Applying VSTs to Tracks	14
5.1.3	Supplementing Dataset with Real Music	14
5.2	Deep Learning Framework	14
5.3	Visualisation	15
5.3.1	Melspectrogram	15
5.3.2	Signal Impulse Response	15
5.4	Audio I/O	15
6	Results	17
6.1	Amplitude Based Effects	17
6.1.1	Distortion	17
6.1.2	Fuzz	17
6.1.3	Amplifier + Cab Simulation	17
6.2	Frequency Based Effects	20
6.2.1	Pitch Shifting (Octave Shift)	20
6.3	Time Based Effects	21
6.3.1	Chorus	21
6.3.2	A Note on Delay and Reverb	21
7	Conclusion	24
7.1	Amplitude Based Effects	24
7.2	Frequency Based Effects	24
7.3	Time Based Effects	24
7.4	Final Thoughts	25

Chapter 1

Introduction

1.1 Musical Effects

Musical effects are transformations which can be applied to an audio track to change the sound in some way [1]. These can take the form of simple tweaks to the frequency range of the track to applying pitch modulation and beyond.

Musical effects can completely change the way a track sounds and the art of mixing a track is - in many ways - more complicated than the composition and performance of the track. Sound engineers have a huge range of choices and responsibility when it comes to getting the soundscape correct for the final product.

The music industry is one of the few domains which is still to fully embrace digital technology. Many recording studios and musicians still make use of analogue equipment such as tape, vinyl, and vacuum tubes [1]. As these technologies become more outdated and niche, the cost of maintaining and replacing equipment rises. As such, it is essential that software modelling catches up to the performance and quality of these older technologies.

1.2 Musical Effect Modelling

The main standard for producing digital effects for use with audio is through the Virtual Studio Technology (VST) protocol. This protocol defines a standardised definition for how a VST host would behave, allowing VST Plugins to be created which manipulate or use the audio track in some way.

The technology itself is fairly flexible and allows everything from audio visualisation tools to CPU intensive transformations to be applied. These effects can also be stacked and ordered easily and intuitively.

Most effects however, are limited to digital signal processing (DSP) based transformations, the likes of which are unable to replicate the non-linearity of many analogue methods. This is cause for concern in the music industry and is the reason many professionals and hobbyists claim that they are inadequate for anything more than practice and experimentation.

1.3 Machine Learning

In recent years, a huge amount of research has gone into machine learning. Its application frequently makes news for revolutionising the various industries it is applied to. As such, it would be interesting to see the effectiveness of machine learning in this field.

Of particular interest to us were convolutional neural networks and Long-Short Term Memory (LSTM) networks, which both seem to be at the forefront of deep learning research. However, the application of machine learning to audio data is in its infancy, resulting in quite a challenging implementation with many layers to it and several dependencies.

The idea of the project is to create two audio tracks, one which is a clean signal from an instrument and one which is the result of applying a VST effect to that clean track. We would then create a neural network which would take samples from the clean signal and be trained to produce the value of the corresponding sample in the processed track.

Chapter 2

Background

2.1 Nebula VST

Nebula yields the best results current modelling technologies can achieve. It achieves this using Volterra Kernel Modelling to better represent non-linear behaviours and exhibiting some level of memory capacity. This allows it to model time-based effects and non-linearity more effectively than primitive DSP approaches.

In doing this, it has become one of the most praised pieces of modelling software, criticised only for its system requirements and its steep learning curve. Requiring a minimum of 8GB of RAM with a recommended amount of 16-128GB, it is a rather demanding piece of software.

Given that Nebula is the best plugin available at this time, for this project to be considered a success, it should be as good as or better than Nebula.

2.2 Project Magenta

Project Magenta is a venture by the Google Brain Team, their goal is to explore the use of machine learning in creating art, with a particular interest in music. Most of their research is in audio synthesis with some ventures into genre classification and similar problems.

However, this project does raise some interesting questions about the ability of deep learning to characterise the complexities of music.

Most of their work however, was limited to interacting with MIDI data rather than with raw audio samples, perhaps limiting the complexity of the problem, making it an unfair comparison in terms of complexity to effect modelling.

2.3 WaveNet

Another use of deep learning for audio synthesis is in the field of generating realistic text-to-speech. DeepMind's WaveNet is a generative network, trained using raw audio samples to predict the value of the next sample in the series.

The project found that while generally an LSTM is better suited to this kind of time-series problem, they could achieve similar results using stacked convolutional layers applying dilated casual convolutions. At each layer they would double the dilation amount up to a maximum before starting again at 1. This increases the receptive field of the network in a similar way to how an LSTM selectively remembers traits while maintaining the ease of training of a convolutional network.

Another way in which the WaveNet implementation differs from a standard convolutional network is in its output layer. Instead of predicting a sample as a floating point value representing the amplitude, the output is a one-hot encoded vector where the largest index corresponds with the μ -Law encoded integer of the sample.

$$F(x) = \text{sgn}(x) \frac{\ln(1+\mu|x|)}{\ln(1+\mu)} - 1 \leq x \leq 1$$

μ Law as expressed in Cisco, 2006

This network has proven itself to be capable of generating extremely effective, almost human sounding text-to-speech and is far easier to train than a standard LSTM-based approach.

Chapter 3

Aims

3.1 Effect Choices

The aim of this project is to explore the viability and success of deep learning in effect modelling at a fairly general level. As such, a multitude of different effects should be tested.

While there is a huge range of commonly used effects, and an even larger range of different variations of these same effects with subtly different responses, most of them fit in to one of three groups.

Amplitude Based	Frequency Based	Time Based
Distortion Fuzz Compression Amplifier/Cab Simulation	High/Low Pass Filter Pitch Shift/Octave	Chorus Delay Reverb

Narrowing the scope of the project to include only these effects results in a good indication of viability across a multitude of effects, while keeping the project scope at an attainable level.

3.2 Deep Learning

Deep learning has demonstrated extreme potential in the realm of text-to-speech and audio generation. As such, it is possible that it would prove similarly effective in effect modelling.

A few different types of deep learning were of particular interest. With regards to amplitude based effects, most of the transformation is carried out across a narrow time-frame, and non-linearity is limited. As such, for these effects it was hypothesised that a simple Convolutional Net with a window of narrow window of samples as the input would be sufficient. The same was thought for the frequency based effects such as filters.

For Time Based effects however, the active window of the effect can be much larger, with reverb sometimes being used with multiple second long tail. Given that music is generally sampled at 44,100Hz, this means that for the network to model this, it would require more than 44,100 samples as its input vector. This is obviously impractical as the memory requirements and time to train the network would be impractical.

In order to try and work around this, LSTM networks would be explored as a potential remedy. These would be able to learn which aspects of the track need to be remembered to best replicate the effect.

Chapter 4

Methods

4.1 Network Architecture

4.1.1 Convolutional Networks

Apart from WaveNet, guidance for creating a network architecture suitable for capturing the complexities of audio data were hard to come by. Most of the experience I had with machine learning was limited to completing example problems and tutorials.

As such, the first goal was to adapt one of these example networks to receive audio data rather than text or image data which the tutorial used. The network I chose to adapt was one which performed well on the MNIST image classification dataset.

This network consisted of 2 Convolutional Layers with Rectified Linear Unit (ReLU) activation layers between, the results of which were then forwarded into a Max Pooling layer before a series of fully connected layers. A simplified diagram of this is shown in Figure 4.1

4.1.2 Long-Short Term Networks

It was thought that while a convolutional network would be impractical in learning the spatial elements required to reproduce a time-based effect such as reverb or chorus. Some kind of recursive neural network (RNN) was hypothesised to be able to replicate this behaviour as it would take the samples from the track in one at a time and learn what properties to remember and for how long.

The downside to this was that training time is much slower as the network has a great deal more properties to learn than a typical feed-forward network.

4.1.3 WaveNet

Another idea was to implement the same basic network as WaveNet. Using dilated casual convolutions to classify the predicted value based on a one-hot encoded vector. This worked extremely well for generating speech data, exhibiting a “subjective naturalness never before seen.”



Figure 4.1: A Simplified Diagram of the Convolutional Network used

An implementation of WaveNet was used in this project but proved impractical as One-Hot Encoding/Decoding of 24 bit audio is slow due to the linear time of decoding. While this worked fine for 8 bit speech data, such a sample depth proved too low to represent audio of a musical nature with enough accuracy to be considered successful.

4.2 Evaluating Success

4.2.1 Melspectograms

Spectograms are a graphical representation of an audio signal. They show detailed information about the frequencies present in a signal over time. The resulting representation is often used in speech recognition problems as a means of feature extraction. Allowing the system to identify the words spoken. As such, they can be used as a quantitative means of evaluating how closely the network output matches the VST output.

The process for creating Melspectograms is based on mapping the signal to Mel Scale using the Fourier Transform. This can then be represented as a heatmap of frequencies in Hz which are present in the audio. Figure 4.2 shows an example of this on an audio track.

This is a vital tool in determining how effectively frequency based effects are being modelled as it shows any changes to the EQ of the track. For a model to be considered successful, the spectogram from the model should look more similar to the spectogram of the VST than the clean signal. This demonstrates that the EQ profile of the model output is exhibiting the same changes that the VST made.

4.2.2 Impulse Response

Impulse responses are used to represent distinct audible events. They can be visualised by simple graphing the amplitudes over time of a sample. This generates a line showing the intensity of the signal over a given time period.

Impulse responses can be useful in identifying the perceived volume or gain applied to a signal, making it a vital tool in determining how well the network is modelling distortion and other amplitude based effects.

Figure 4.3 shows an example of an impulse response graph.

In the interest of ease of comparison, in the results section the impulse response of the clean signal, the VST signal, and the model's effort will be plotted in on the same graph for comparison (as shown in figure 5.1). This allows the three signals to be compared easily. For the model to be considered successful, its curve should be more similar to that of the VST effect than the Clean signal. This would mean that any changes to amplitude the VST makes, the model is attempting to capture with some degree of success.

4.2.3 ABX Testing

ABX testing takes the form of having 3 tracks, A which is our clean signal, B which is our processed effect, and X which is the output of the network. The listener is not informed which track is which and has to say whether X sounds more like A or B.




Figure 4.2: A Melspectrogram of a guitar recording

For our project to be considered successful, the results would show extreme bias in X sounding more like B than A. As this means that the network is producing tracks which sound more like the processed effect than the clean input.




Figure 4.3: An Impulse Response Graph of 300 samples in a track

Chapter 5

Implementation

5.1 Generating Test Data

In any deep learning scenario, one of the most important factors in success is in having a large and varied dataset for training. As such, for this project, it was decided that we would generate random audio data as needed, giving a theoretically infinite amount of training data.

While random data ensures a significantly large dataset, the number of parameters which could be randomised was limited. In case this resulted in overfitting, it was also decided to supplement the randomised data with real music.

5.1.1 Random Audio Signals

The random audio data was generated using the `scipy.signal` module. The function `chirp` was used to generate an audio signal which rises or falls in pitch from one frequency to another. The frequencies are both randomised from the range of 80Hz to 1200Hz as this is roughly the frequency range of an electric guitar. The speed at which this frequency sweep occurs is also randomised. Frequency sweeping was thought to imitate a guitar string being bent or slid from one note to another.

Numpy was also used to generate sine waves at fixed frequencies from the above range. This would simulate a single sustained note from a guitar or other instrument.

To further add variety, `scipy`'s `square` and `sawtooth` functions were used with random frequencies to further add to the variety of different sounds that can appear in the resulting audio tracks.

Beyond this, the generated data is split into 'segments' which was done to replicate the way that music can vary in terms of speed, intensity and general level of complexity. Each segment is 10 seconds long.

This means that every 10 seconds of audio data generated varies in terms of how many different notes or waveforms are being expressed at a given time, and how big the gaps between notes are. This was introduced to replicate the fact that there is rarely just a signal note playing in a musical piece. This could be important for modelling the time-based effects as it adds a greater variety to the audio which may need to be remembered to successfully model chorus, delay and reverb.

Generating the audio was carried out in a Python script called `generate.py` which two integer arguments, the first determines how many audio tracks are to be generated and the second determines the number of segments

per track.

5.1.2 Applying VSTs to Tracks

To train the network, we need a clean signal which takes the form of the output from `generate.py`, we then also need files which represent that track after it has been processed by an effect.

In order to achieve this, a VST host which could be invoked through the command line was preferred. Initially a command line utility known as MrsWatson was selected for this. It could be invoked with a command such as `mrswatson --input mysong.wav --output out.wav --plugin myplugin` to apply an effect called `myplugin` to `mysong.wav`.

However, MrsWatson proved to behave inconsistently. Using it with Linux required the plugins to have been specifically compiled for Linux and since there are so few audio professionals working on this platform the number of effects available were seriously limited. Beyond this, the quality of the output was poor, often exhibiting hiss or other signal noise being introduced.

As such, a new solution was required. This came in the form of Reaper. A fully-featured Digital Audio Workstation. While Reaper has a huge range of features, the feature of interest in this project is its batch mode which can be invoked with the command line. This mode allows a list of files to be given in a line-delimited file and a Reaper signal chain file to be given to be applied to each of these files.

A bash script was created to automatically generate the files using `generate.py`, the dataset would be saved in a folder called 'dataset' The script would then run Reapers batch mode on the files putting the resulting files in 'dataset/processed' with the same file name as their corresponding source files.

5.1.3 Supplementing Dataset with Real Music

The script was later adapted to incorporate any audio files placed in the 'training_audio' folder into this process. Allowing users to easily supplement the random training data with real music. This could potentially solve issues with overfitting and make it easier for the network to learn more subtle effect features.

5.2 Deep Learning Framework

After considering the various options for machine learning in Python, the framework of choice was PyTorch. PyTorch facilitates a tight integration between the Python code which defines the network architecture and the C/C++ code which runs the neural network defined in the Python code. Using a framework like Caffe or TensorFlow, the state of the network is hidden from the Python runtime as it exists as a separate process. PyTorch works around this with only a small impact on speed and memory efficiency.

This allows for easy debugging as the Python debugger and simple `print` statements can be used to verify the behaviour of the network at a given time. This alleviates the learning curve of having to use something like TensorBoard for TensorFlow to visualise the network state.

Beyond this, PyTorch has a very 'Python-first' philosophy, allowing Tensors to be manipulated using standard Python list slicing, it's even possible to perform list comprehensions on them.

PyTorch has three main modules which were used in this project:

- The base module `torch`, which provides access to the `torch.Tensor` and its various subclasses.
- `torch.nn` which provides the necessary implementations of neural network layers.
 - Also provides `nn.Sequential`, which can be used to stack various layers into one class for ease of use and to eliminate the boilerplate code which often is required of other frameworks.
- `torch.autograd` which provides the `Variable` class, this encapsulates a `Tensor` and allows it to be used with backward propagation for Deep Learning to be carried out.

5.3 Visualisation

5.3.1 Melspectrogram

This was achieved using `Matplotlib` and `LibRosa`. `LibRosa` has a range of functions for visualising audio data, including the function `melspectrogram` found within the `librosa.feature` module. This generates a Numpy array. The `specshow` method in `librosa.display` acts as a wrapper for `matplotlib.pyplot` and generates the spectrogram when given the numpy array generated by `melspectrogram`.

5.3.2 Signal Impulse Response

Impulse response graphs are relatively simple to implement from the data we get from `librosa.core.load` which is used to read in audio files. The result is an array of amplitudes over time which can simply be plotted in a line using `matplotlib.pyplot`.

We can better compare the three audio files by plotting them the three of them on the same axis as is shown in figure 5.1

5.4 Audio I/O

The project utilises `LibRosa` for reading and writing `wav` files. This was chosen as it supports a multitude of different output formats. Early versions of the project were using `SciPy.IO.wavfile` for this but it was limited to 64 bit output files, which was unnecessarily large for audio data and required too much storage space.

`LibRosa` will output the file with the smallest sample depth required to adequately represent the track, creating much smaller files. It is also much more robust and can even read `MP3`, for example if users want to include their own audio into the training data without having to convert it to `wav`.




Figure 5.1: Demonstrates the Clean Signal (Orange), the VST Signal (Green), and the Model Signal (Blue). We can see the model is generating a signal which lies somewhere between the clean signal and the VST, this shows the model is learning the necessary transformation.

Chapter 6

Results

For a model to be considered successful, the Spectrogram and Impulse responses should look more similar to the VST output than the clean signal.

6.1 Amplitude Based Effects

6.1.1 Distortion

This model was created using a Convolutional Network consisting of 6 convolutional layers, which have a ReLU activation layer between them, before a Max Pooling Layer and a Fully Connected section which consists of 3 Linear layers.

As demonstrated in the figures 6.2, 6.3, and 6.4, the distortion unit adds a huge amount to the intensity of all frequencies, with a bias toward the mid ranges. This behaviour is captured fairly well by the model. Figure 6.1 also demonstrates that the amplitude boost is also learned by the model.

6.1.2 Fuzz

Fuzz seems to be modelled with similar success to distortion, which makes sense since it belongs to the same family of effects. Figures 6.5, 6.6, 6.7, and 6.8 demonstrate that the model is able to capture the characteristics of the fuzz effect in terms of both EQ range and amplitude.

6.1.3 Amplifier + Cab Simulation

Amp Simulation is a historically difficult task. Modelling amplifiers still have a tiny market share and are some of the most criticised applications of software modelling. This is reflected in the fact that the model completely fails to learn the required transformation to mimic amp simulation.

While the spectrogram (Figures 6.10, 6.11, and 6.12) seems to confirm that some aspects of the EQ profile are learned, it fails to match as closely as the previous effects. Beyond this, the VST seems to add a lot more gain, as demonstrated in 6.9. This increased amplitude is expressed as the audio track becoming clipped and distorted in ways that are not musical or intentional.




Figure 6.1: Distortion after 30000 iterations over the dataset




Figure 6.2: Clean Spectrogram (Distortion)




Figure 6.3: Model Spectrogram (Distortion)




Figure 6.4: VST Spectrogram (Distortion)




Figure 6.5: Fuzz after 30000 iterations over the dataset




Figure 6.6: Clean Spectrogram (Fuzz)




Figure 6.7: Model Spectrogram (Fuzz)




Figure 6.8: VST Spectrogram (Fuzz)




Figure 6.9: Amp Simulation after 30000 iterations over the dataset




Figure 6.10: Clean Spectrogram (Amp Sim)

This is perhaps due to the fact that there are a lot more behaviours to learn in an amp simulation than there is with a basic distortion effect, in fact amp simulation may even bridge the two categories of Amplitude based effects and Frequency based effects since it operates over both of these domains.

6.2 Frequency Based Effects

6.2.1 Pitch Shifting (Octave Shift)

Pitch shifting is an interesting effect, the domain of the effect is solely in the EQ profile, unlike amp simulation which spans both domains. As such, it is useful in confirming whether the model is capable of learning these EQ based transformations.

Figures 6.14, 6.16, and 6.15 seem to confirm that some of this EQ based behaviour is captured and replicated. Listening to the track also confirms this to be true. However, as shown in 6.13, the model adds a lot of additional noise to the signal, this comes across in the track as hiss and clipping. The model output is significantly louder




Figure 6.11: Model Spectrogram (Amp Sim)




Figure 6.12: VST Spectrogram (Amp Sim)

than both the clean signal and the VST processed track.

6.3 Time Based Effects

6.3.1 Chorus

Time based effects were by far the least well performing category in the project. Early tests with the standard convolutional network used a huge sample size of around 5000. This was the largest input into the network that could be achieved with 16GB of RAM. Unfortunately, this failed to characterise the effect.

In order to combat this, an LSTM network was used instead. This took the form of a 7 layer LSTM with a hidden size of 800. A new sampler, the `randomSequentialSampler` was created to train for X seconds of audio in order before switching to an entirely new section of the track. This was to combat overfitting that comes from training on a track sequentially. Sampling was completely random for other effects. Random sampling would not have worked on the LSTM as it would have been impossible for the network to learn what to remember if it was in random order.

Despite this, the LSTM failed to characterise the behaviour of the VST and in fact did not produce an audible track at all, even after 2 hours of training. The time to train these networks is slow and even with different numbers of layers and hidden sizes the results were the same.

From figure 6.17, we can see that while there is a variation to the signal in both the Clean and VST output, the LSTM seems to be outputting either very similar or the same value for each sample. It is unclear what would cause such a behaviour. This results in no audible signal being present in the file.

6.3.2 A Note on Delay and Reverb

The results of these effects have been omitted from the paper due to space constraints and the fact that they produced the exact same output as the Chorus effect. This is potentially due to a limitation with the LSTM network in terms of training time and system resources, a network sophisticated enough to model these effects can not be created on my system.




Figure 6.13: Octave Pitch Shift after 30000 iterations over the dataset




Figure 6.14: Clean Spectrogram (Pitch Shift)




Figure 6.15: Model Spectrogram (Pitch Shift)




Figure 6.16: VST Spectrogram (Pitch Shift)




Figure 6.17: Chorus after 30000 iterations over the dataset




Figure 6.18: Clean Spectrogram (Chorus)




Figure 6.19: Model Spectrogram (Chorus)




Figure 6.20: VST Spectrogram (Chorus)

Chapter 7

Conclusion

7.1 Amplitude Based Effects

While the characteristics are modelled with limited success, the tracks generated feature artifacts such as introducing additional hiss or static noise into the signal. This prevents them from being usable as pieces of music and as such makes it difficult to justify the project as anything more than a proof of concept.

7.2 Frequency Based Effects

The network is able to learn some of the transformations to the EQ profile of the track but the output is unconvincing at best. The audio clips and there are clear artifacts in the signal which ruin what little success is reached in this domain. Perhaps a more sophisticated network would have better success in this domain, however the architecture of which has so far been undiscovered.

7.3 Time Based Effects

The LSTM Network failed to produce audio which exhibited the qualities of Chorus, Reverb, or Delay. This is perhaps due to the simple fact that with audio being sampled at 44,100 samples per second, and the active window of these effects potentially spanning more than 0.5 seconds, that the network is simply unable to train itself on what to remember.

Perhaps the WaveNet solution would have had more success with this, however the implementation of it proved impractical for audio data. Given that one-hot decoding is in linear time based on the sample depth of the audio track. This works fine for speech data which is typically 8 bit but for musical data, we require a sample depth of 24 or greater to express the required complexity. As such, this process of one-hot decoding adds an inordinate amount of training time to the point where the network failed to train anything meaningful despite being left overnight.

7.4 Final Thoughts

This project shows potential for deep learning to be used in musical effect modelling. Despite not reaching a production ready solution, with the number of hyperparameters and different network architectures available, there must exist a network which will succeed where this project failed.

Results are therefore a proof of concept rather than a final verdict on the viability of this technique in modelling.